

Medical Imaging Visualization Lab Report

Carles Matoses Gimenez¹

¹Universitat Politècnica de Catalunya - BarcelonaTech

Imanol Munoz-Pandiella

Abstract—This report contains information about the real time volumetric rendering shader. It explains how DVR works with the ray marching algorithm and front to back compositing. It also explains the different techniques used to improve the quality of the rendering, such as lighting, transfer functions and optimizations like empty space skipping and early ray termination. Finally, it shows some results and conclusions about the implementation.

Keywords— $\text{\LaTeX}$ class, lab report, academic paper, tau-book class, volumetric rendering.

1. Introduction

In this lab report, we present the implementation of a real-time volume rendering shader using Direct Volume Rendering (DVR) techniques. The shader is designed to visualize 3D volumetric data by employing ray marching and front-to-back compositing methods. Additionally, we incorporate fresnel effect to enhance certain visual aspects of the rendering.

2. Rendering

We are provided a **.raw file containing a row** (or column) of **scalars**. The name of the file contains the dimensions of the volume. We create a **3D texture in WebGL** to sample the volume data. The vertex shader will receive the vertices, normals and texture of a cube as input. For convenience, we will output the vertex position in world space. In the fragment shader, we will receive the 3D texture, light attributes, positions in world space and TF parameters as uniforms.

```
1 vec3 camera;
2 vec3 ray_o = camera;
3 vec3 ray_d = normalize(frag - ray_o);
4 vec3 d = ray_d * step_size(ray_d);
5
6 vec3 p_in, p_out = intersect_box(ray_o, ray_d);
7 p_in += jitter;
8
9 vec3 C = vec3(0.0);
10 float A = 0.0;
11 while rayMarching:
12     vec3 p = p_in + i*d;
13     // Beer-Lambert
14     float alpha = sample_opacity(p, d);
15     if alpha < threshold: continue;
16
17     vec3 normal = -gradient(p, d);
18     vec3 f_phong = phong(p, normal, l, camera, d);
19     float tr = 1.0 - acc_a(p, l, step_size(l-p));
20     f_color = f_phong * tr;
21
22     A += (1.0-A) * alpha;
23     C += (1.0-A) * alpha * f_color;
24
25     if (A >= 0.95) break;
26     i++
27
28 C = pow(C, vec3(1.0/2.2));
29 frag_color = vec4(C, 1.0);
```

Code 1. Pseudocode for DVR

The Direct Volume Rendering could be done in view space to avoid some computations, but **we will do it in world space** for clarity. We start by computing the ray direction from the camera position to the fragment position 1.

We assume volume of size -1 to 1 in all axis, this can be extended in the future to support arbitrary axis-aligned bounding boxes. The problem is that we do not have any way to know the physical dimensions of the object before it was captured. The best guess would be to assume it was sampled with the same step size in all axis, but this is not always true.

$$\vec{r}_i = \vec{c} - \vec{f} \quad (1)$$

2.1. Ray-Box Intersection

The ray-box intersection is **computed using the slab method** `intersect_box(origin, direction)`. For each axis, we compute the minimum and maximum intersection distances $t_a^{\min}$ and $t_a^{\max}$ using the parametric equation of the ray. Then, we compute $t_{1,a}$ and $t_{2,a}$ as the minimum and maximum of $t_a^{\min}$ and $t_a^{\max}$. Finally, we compute the near and far intersection distances t_{near} and t_{far} as the maximum of $t_{1,a}$ and the minimum of $t_{2,a}$ respectively. If t_{near} is greater than t_{far} , there is no intersection.

2.2. Sampling Step

For the sampling step, we use `step_size(vec3 ray_direction, vec3 dimensions)` that calculates the step size based on the segment we are traversing. The volume is not required to have same voxel size in all axis, we dynamically **calculate the step size for each ray individually based on the ray direction and the volume dimensions**. This method is meant to optimize ray marching to the minimum required steps at any view direction. We use the Nyquist-Shannon theorem to avoid aliasing, so we set the step size to half of the minimum distance between voxel centers along the ray direction. This is not done in the function for reutilization purposes.

2.3. Transfer Function

The **transfer function** `transfer_function_sigma(float density, float step_size)` maps the density value to opacity. Color will be implemented in the lighting function. The TF **works as a rectangular window function (also Tukey window)**, where values below a minimum density are fully transparent and values above a maximum density are fully opaque. We add two more controllers to create a smooth transition between the minimum and maximum density values. The implementation is **done with conditionals** to detect the different zones and apply the corresponding formula.

2.4. Beer-Lambert law implementation

Step size opacity independence is achieved with Beer-Lambert law $T = e^{-\sigma \cdot d}$ where T is transmittance, σ is ex-

tion coefficient and d is distance (step size). Sigma is $\sigma = uTFOpacity * (1.0/d)$, the d terms cancel out. The opacity becomes independent of step size and only depends on uTFOpacity.




Figure 1. Visualization showing step size independence using Beer-Lambert law normalization. If step size is not taken into account, we darken the volume faster than expected when using smaller step sizes, leading to inconsistent rendering results.

2.5. Color Accumulation

Color accumulation is performed in `accumulate_color_and_opacity(vec3 p_in, vec3 p_out, float step_size)` using front-to-back ray marching. The ray is sampled uniformly from the near intersection point p_{in} to the far intersection point p_{out} . At each step, the volume density is converted to an extinction coefficient σ via the transfer function.

Opacity is computed using the Beer-Lambert law over the current step length d :

$$\alpha = 1 - \exp(-\sigma \cdot d) \quad (2)$$

If the sample opacity exceeds a small threshold, gradients are evaluated to estimate the local normal, and Phong lighting is applied using `computeLighting`. The lighting model combines ambient, diffuse, and specular terms:

$$\mathbf{C}_{sample} = \mathbf{C}_{ambient} + \mathbf{C}_{diffuse} + \mathbf{C}_{specular} \quad (3)$$

with standard Phong definitions for each component. The shininess parameter α is controlled through the UI (light radius).

To account for light absorption between the current sample position and the light source, the shaded color is multiplied by a transmittance term computed along the shadow ray using `accumulate_opacity`. This secondary ray march calculates the step size based on the light direction. For optimization reasons we do not reduce the step size by half when calculating shadows (becomes unnecessary expensive for non noticeable improvements).

Color and opacity are accumulated using **premultiplied alpha compositing** in a front-to-back manner:

$$\mathbf{C}_{acc} \leftarrow \mathbf{C}_{acc} + (1 - \alpha_{acc}) \cdot (\mathbf{C}_{sample} \cdot \alpha) \quad (4)$$

$$\alpha_{acc} \leftarrow \alpha_{acc} + (1 - \alpha_{acc}) \cdot \alpha \quad (5)$$

This formulation corresponds to a Riemann sum approximation of the volume rendering integral and matches the Beer-Lambert extinction model used in the implementation.

For efficiency, **early ray termination** is applied: once the accumulated opacity exceeds a threshold (e.g. $\alpha_{acc} > 0.95$), the ray marching loop is terminated since further contributions become negligible.

2.6. Gradient Computation for Lighting

We have tried **two methods to compute the gradients**, both with advantages and disadvantages. The **first method uses density values**. This means we do not take the TFs into account. This is perfect for voxels in the margins so they **can show normals even if the TF returns maximum opacity in a rectangular window**. This gives the best results for most cases. **Hollow objects will have inverted normals if looked from the inside**. If TF removes bones and muscles to only leave skin, we will not be able to illuminate it with the current diffuse lightning since normal directions will point outwards as expected. The **second approach is to use opacity values** from the TF. We solve the illumination of hole objects but **we loose density information** when calculating the gradient.




Figure 2. Comparison of normals computed with opacity (left) and density (right). Notice how the opacity normals fail to capture the surface details due to the transfer function removing important density information. Density normals fail to provide correct illumination for hollow objects when viewed from the inside (normals are inverted).

The gradient ∇f at point $\mathbf{p}$ is **computed using central differences**. We sample the volume at six neighboring points displaced by $\Delta s \cdot c_{smooth}$, where Δs is half the step size and c_{smooth} is a smoothing coefficient:

$$\mathbf{p}_l = \mathbf{p} + (-\Delta s \cdot c_{smooth}, 0, 0) \quad (6)$$

$$\mathbf{p}_r = \mathbf{p} + (\Delta s \cdot c_{smooth}, 0, 0) \quad (7)$$

$$\mathbf{p}_b = \mathbf{p} + (0, -\Delta s \cdot c_{smooth}, 0) \quad (8)$$

$$\mathbf{p}_t = \mathbf{p} + (0, \Delta s \cdot c_{smooth}, 0) \quad (9)$$

$$\mathbf{p}_{close} = \mathbf{p} + (0, 0, -\Delta s \cdot c_{smooth}) \quad (10)$$

$$\mathbf{p}_{far} = \mathbf{p} + (0, 0, \Delta s \cdot c_{smooth}) \quad (11)$$

The gradient is then computed as:

$$\nabla f(\mathbf{p}) = \begin{pmatrix} f(\mathbf{p}_r) - f(\mathbf{p}_l) \\ f(\mathbf{p}_t) - f(\mathbf{p}_b) \\ f(\mathbf{p}_{\text{far}}) - f(\mathbf{p}_{\text{close}}) \end{pmatrix} \quad (12)$$

where $f(\mathbf{p})$ represents the sampled density (or opacity) value at position $\mathbf{p}$. The gradient vector points in the direction of maximum density increase. For surface normals, we negate this vector to obtain the outward-pointing normal:

$$\mathbf{n} = -\frac{\nabla f(\mathbf{p})}{|\nabla f(\mathbf{p})|}.$$

Additional problems to normal calculations: **If using density, we need to know the minimum possible value** of density since "void" will take that value (or smaller) for sampling out of the volume points. If done correctly, the model should have values between 0 and infinity but this is not the case in the medical images used in the warmup (from -1024 to 1637).

3. Silhouette

4. Performance

References